

Lecture Outline

Generative
modellingEnergy-based
modelsScore-based
models

Sampling

Computational Statistics & Machine Learning

Lecture 14

Energy-based and Score-based Modelling

Mark Girolami

`mag92@cam.ac.uk`

Department of Engineering

University of Cambridge

October 15, 2025

Overview

M.Girolami

Lecture Outline

- ▶ Generative modelling
- ▶ Energy-based models
- ▶ Score-based models
- ▶ Sampling

Generative
modelling

Energy-based
models

Score-based
models

Sampling

- ▶ Learning models of the data distribution is a significant problem in data science and machine learning
- ▶ Given a dataset $\mathcal{D} = \{x_i\}_{i=1}^N$
- ▶ Assumed generated according to some *unknown* data distribution $x_i \sim p_{\text{data}}(x)$
- ▶ Aim to approximate $p_{\text{data}}(x)$ with parameterised $p_\theta(x)$
- ▶ Learn parameters θ from the data \mathcal{D}


We will cover different methods for learning $p_\theta(x)$, and generating samples $x' \sim p_\theta(x)$

Generative modelling of images

DeepDream (2014)



GANs (2015)



Stable Diffusion, DALL-E, Midjourney (Current)



Lecture Outline

Generative
modelling

Energy-based
models

Score-based
models

Sampling

- ▶ Inspired by statistical physics, where probability of a system state is modelled via an **energy function** $E(x)$
- ▶ The probability of a state x , is determined by the energy function $E(\cdot)$, and the (inverse) temperature β

$$p(x) = \frac{1}{Z_\beta} \exp(-\beta E(x))$$

- ▶ Known as the Boltzmann or Gibbs distribution
- ▶ Normalisation constant

$$Z_\beta = \int \exp(-\beta E(x)) dx$$

is the ‘partition function’

- ▶ Partition function generally difficult to compute - requires integration over all states

- ▶ EBMs directly parameterise the energy function:

$$p_\theta(x) \propto \exp(-E_\theta(x))$$

- ▶ Boltzmann machines were introduced by Geoffrey Hinton and Terry Sejnowski in 1985
- ▶ Early example of neural networks modelling complex distributions
- ▶ Trained via **maximum-likelihood** or **contrastive divergence**

$$p_\theta(x) = \frac{1}{Z_\theta} \exp(-E_\theta(x)), \quad Z_\theta = \int \exp(-E_\theta(x)) dx$$

► Maximum-likelihood

$$\mathcal{L}(\theta) := \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log p_\theta(x)], \quad \theta^* = \underset{\theta}{\operatorname{argmax}} \{\mathcal{L}(\theta)\}$$

► Parameters optimised via gradient methods

$$\nabla_\theta \log p_\theta(x) = -\nabla_\theta E_\theta(x) - \nabla_\theta \log Z_\theta$$

► Second term involving partition function Z_θ can be shown to be equivalent to

$$\nabla_\theta \log Z_\theta = -\mathbb{E}_{x \sim p_\theta(x)} [\nabla_\theta E_\theta(x)]$$

► Giving the gradient of the objective as

$$\nabla_\theta \mathcal{L}(\theta) = -\mathbb{E}_{x \sim p_{\text{data}}(x)} [\nabla_\theta E_\theta(x)] + \mathbb{E}_{x \sim p_\theta(x)} [\nabla_\theta E_\theta(x)]$$

Contrastive Divergence

The maximum likelihood gradient requires sampling from the model p_θ (second term).

- ▶ Generally requires using expensive MCMC methods
- ▶ **Contrastive divergence** uses practical (but biased) approximation:

1. For each x_i , initialise MCMC chain (p_θ) at $x_i^{(0)} := x_i$
2. Run for short sequence of k steps, giving $x_i^{(0)}, \dots, x_i^{(k)}$
3. Compute gradient estimate

$$\hat{\mathcal{L}} = -\frac{1}{N} \sum_{i=1}^N \nabla_\theta E_\theta(x_i^{(0)}) + \frac{1}{N} \sum_{i=1}^N \nabla_\theta E_\theta(x_i^{(k)})$$

4. Update parameters $\theta \leftarrow \theta + \eta \hat{\mathcal{L}}$
(or other gradient-based method)
- ▶ Downside is CD is biased, and MCMC can be very slow.
Due to computing this partition function

Score-based models

Score-based models, instead of directly modelling the energy function, aim to match the **score** of the data-distribution

- ▶ Denoting the score function $s_\theta(x) := \nabla_x \log p_\theta(x)$
- ▶ Working with the score means we don't have to calculate the partition function Z_θ :

$$s_\theta(x) = -\nabla_x E_\theta(x) - \underbrace{\nabla_x \log Z_\theta}_{=0} = -\nabla_x E_\theta(x) \quad (1)$$

- ▶ Score matching, introduced by Hyvärinen (2005), can be used to train these score models
- ▶ Aim to minimise the Fisher divergence between the true and approximate score:

$$J(\theta) = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}(x)} \left[\|\nabla_x \log p_{\text{data}}(x) - s_\theta(x)\|^2 \right]$$

as we are only given samples from p_{data} , this is intractable in its current form - can be rearranged

Score matching

Fisher divergence can be shown to be equivalent to the following:

$$J_{SM}(\theta) = \mathbb{E}_{x \sim p_{\text{data}}(x)} \left[\text{tr}(\nabla_x s_\theta(x)) + \frac{1}{2} \|s_\theta(x)\|^2 \right]$$

which can be estimated using data samples via MC.

- ▶ However... the trace of the Jacobian, $\text{tr}(\nabla_x s_\theta(x))$, is expensive for high-dimensions.

Sliced score matching uses random projections to approximate this trace.

$$\text{tr}(\nabla_x s_\theta(x)) = \mathbb{E}_{v \sim p(v)} [v^\top \nabla_x s_\theta(x) v]$$

- ▶ Applying this gives the objective function for sliced score matching:

$$J_{\text{SSM}}(\theta) = \mathbb{E}_{v \sim p(v)} \mathbb{E}_{x \sim p_{\text{data}}} \left[v^\top \nabla_x s_\theta(x) v + \frac{1}{2} \|s_\theta(x)\|^2 \right]$$

- ▶ The term $v^\top \nabla_x s_\theta(x) v$ can be computed efficiently in high dimensions using Jacobian-vector products

how are these trained in practice?

Training Score-based Models

For standard score matching, we use the data samples to update parameters using any gradient descent scheme.

- ▶ Initialise θ_0 , then iterate $t = 1, \dots$

$$\hat{J}_{\text{SM}}(\theta_t) = \frac{1}{N} \sum_{i=1}^N \text{tr}(\nabla_x s_{\theta_t}(x_i)) + \frac{1}{2} \|s_{\theta_t}(x_i)\|^2$$

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \hat{J}_{\text{SM}}(\theta_t)$$

- ▶ or stochastic gradient descent, using batch $B_t \subset \{1, \dots, N\}$, size $m = |B_t|$

$$\tilde{J}_{\text{SM}}(\theta_t) = \frac{1}{m} \sum_{i \in B_t} \text{tr}(\nabla_x s_{\theta_t}(x_i)) + \frac{1}{2} \|s_{\theta_t}(x_i)\|^2$$

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \tilde{J}_{\text{SM}}(\theta_t)$$

Training Score-based Models

- ▶ Sliced score matching requires sampling from $p(v)$ to compute the random projection.
- ▶ In practice, typically one sample is used for computing this estimate

$$\hat{J}_{\text{SSM}}(\theta_t) = \frac{1}{N} \sum_{i=1}^N v_i^\top \nabla_x s_{\theta_t}(x_i) v_i + \frac{1}{2} \|s_{\theta_t}(x_i)\|^2$$

where $v_i \sim p(v)$.

- ▶ Once we have estimated parameters θ^* , we can evaluate the estimated score as $s_{\theta^*}(x)$.
- ▶ How can we generate new samples using the estimated score?
- ▶ Can't use MCMC like for EBMs, as we can't evaluate $p_{\theta^*}(x)$, we can only evaluate $s_{\theta^*}(x)$
- ▶ Langevin dynamics

- ▶ Initial sample $x_0 \sim p_0(x)$ from some initial distribution
- ▶ Recall Langevin SDE of the form

$$dx_t = \nabla_x U(x) dt + \sqrt{2} dB_t$$

has invariant density $p(x) \propto \exp(-U(x))$.

- ▶ Setting $\nabla_x U(x) = s_{\theta^*}(x)$ gives us $p_{\theta^*}(x)$ as invariant density
- ▶ Simulate Langevin dynamics, e.g. ULA, with step-size $\gamma > 0$

$$x^{(k+1)} = x^{(k)} + \gamma s_{\theta^*}(x^{(k)}) + \sqrt{2\gamma} z^{(k)}$$

where $z^{(k)} \sim \mathcal{N}(0, I)$.